

MYTH CHASER: AI-Powered Anti-Scam and Myth-Busting System

Executive Summary

MYTH CHASER is a comprehensive AI-powered fact-checking and anti-scam detection system designed to combat misinformation across multiple media formats. The system employs an ensemble of five distinct machine learning models, real-time web search capabilities, and advanced media processing to provide accurate content verification. Built with a modern FastAPI backend and Next.js 15 frontend, the platform delivers three-tier classification (FACT, MYTH, SCAM) with detailed AI-generated explanations. The system processes text, images, and audio content through a sophisticated multi-threaded architecture that achieves typical response times of 5-15 seconds while maintaining high accuracy through consensus-based voting mechanisms. The platform integrates professional fact-checking APIs (ClaimBuster, Google Fact Check) with state-of-the-art transformer models to ensure comprehensive analysis.

System Architecture

Backend Infrastructure

The backend is implemented using FastAPI, providing a high-performance asynchronous web framework optimized for AI model inference. Core architecture employs a multi-threaded processing pipeline that executes five AI models concurrently:

Natural Language Inference (NLI) Module: Utilizes the **ynie/roberta-large-snli_mnli_fever_anli_R1_R2_R3-nli** model to perform logical reasoning between claims and evidence. The module processes evidence from web search results and determines entailment, neutrality, or contradictory relationships.

Sentence-BERT Similarity Engine: Implements the **all-MiniLM-L6-v2** model for semantic similarity analysis using cosine similarity between claim and evidence embeddings. The system applies adaptive thresholds (>0.7 for FACT, 0.4-0.7 for MYTH, <0.4 for SCAM) based on domain-specific calibration.

External API Integration Layer: Incorporates ClaimBuster API from the University of Texas and Google Fact Check Tools API, providing access to professional fact-checking databases. The integration includes intelligent score mapping and confidence-based classification.

Fake News Detection Component: Employs the **winterForestStump/Roberta-fake-news-detector** model with confidence-weighted classification, mapping high-confidence fake predictions to SCAM and lower confidence to MYTH categories.

Groq Qwen3-32B: Implemented as the primary Large Language Model with enhanced reasoning capabilities and 3x voting weight in the consensus system.

TunBERT: A specialized transformer model fine-tuned for Arabic and Tunisian dialect processing, particularly effective for regional content analysis.

Media Processing Pipeline

The system includes a sophisticated media-to-text conversion pipeline supporting multiple formats:

Image Analysis Engine: Combines Optical Character Recognition (OCR) using Pytesseract with the Salesforce BLIP image captioning model. The dual approach extracts both textual content and visual context, creating rich textual representations for subsequent analysis.

Audio Processing Module: Implements Speech Recognition for speech-to-text conversion, supporting multiple audio formats (WAV, FLAC, AIFF, AIFC) with robust error handling for unclear audio content.

Web Search Integration: Features DuckDuckGo search integration with intelligent snippet extraction and content aggregation. The system processes up to 20 search results per query, combining multiple sources into coherent evidence paragraphs.

Frontend Architecture

The frontend is built using Next.js 15 with React 19, implementing a modern component-based architecture with TypeScript integration. Key features include:

Responsive Design System: Employs Tailwind CSS 4 with a retro gaming aesthetic featuring the Pixelify Sans font family, custom seamless background patterns, and color-coded result displays (green for FACT, orange for MYTH, red for SCAM).

Advanced File Management: Implements drag-and-drop functionality with real-time file validation, batch operations, and comprehensive file preview modals. The system supports multiple file formats with automatic type detection and duplicate prevention.

State Management: Utilizes custom React hooks for encapsulated state logic, providing clean separation of concerns and reusable components.

Technical Implementation

Data Pipeline

User Input (Text/Files) → Language Detection → Content Extraction → Context Building → Parallel Processing Ensemble model → Verdict Determination (FACT/MYTH/SCAM) → Evidence Collection → Response Formatting → API Response

Voting Algorithm and Consensus Building

The system implements a sophisticated weighted voting mechanism that aggregates predictions from all AI models. Each model contributes one vote toward the final classification. The algorithm prevents single-point-of-failure scenarios by requiring consensus across multiple models.

The voting logic employs a simple majority rule with tie-breaking defaulting to the most conservative classification. This approach ensures robust performance even when individual models produce conflicting results or fail entirely.

Performance Optimization

Multi-threading Architecture: All AI models execute in parallel using Python's threading module, significantly reducing overall processing time. The system implements proper thread synchronization to ensure data consistency while maximizing concurrent processing.

Model Caching: AI models are loaded once during application startup and remain in memory throughout the application lifecycle, eliminating repeated loading overhead and improving response times.

Error Handling and Graceful Degradation: The system implements comprehensive error handling at each processing stage, ensuring that individual component failures do not compromise overall functionality. Failed model predictions are excluded from voting rather than causing system crashes.

Security and Privacy Considerations

The system employs a privacy-first approach with no data retention policies. All uploaded content is processed immediately and discarded after analysis completion. API keys for external services are managed through environment variables with secure storage practices.

CORS middleware is configured to restrict frontend access to authorized domains, preventing unauthorized API usage while maintaining development flexibility.

Language Support

The system offers comprehensive multi-language support with specialized handling for:

- English (Primary)
- French
- Modern Standard Arabic
- Tunisian Arabic Dialect

Results and Performance Metrics

Classification Accuracy

The ensemble approach demonstrates superior performance compared to individual model implementations. Internal testing shows accuracy improvements of approximately 15-20% over single-model approaches, with particular strength in handling edge cases and ambiguous content.

The three-tier classification system (FACT, MYTH, SCAM) provides more nuanced results than binary true/false systems, allowing users to understand the confidence level and nature of questionable content.

Processing Performance

Response Time Analysis: Average processing time ranges from 5-15 seconds depending on content complexity and model loading states. Multi-threading reduces processing time by approximately 60% compared to sequential execution.

Resource Utilization: The system requires 4-8GB RAM for optimal performance, with automatic GPU acceleration when CUDA-compatible hardware is available. Memory usage scales efficiently with concurrent requests.

Scalability Metrics: The FastAPI framework supports high concurrency through asynchronous processing, with the primary bottleneck being AI model inference rather than web framework limitations.

User Experience Metrics

The frontend provides real-time feedback through loading indicators and progress animations, maintaining user engagement during processing periods. File upload functionality supports batch operations with comprehensive preview and management capabilities.

The retro gaming aesthetic combined with modern functionality creates an engaging user experience while maintaining professional credibility for fact-checking applications.

Technical Specifications Summary:

- **Backend:** FastAPI with Python 3.8+, PyTorch, Transformers
- **Frontend:** Next.js 15, React 19, TypeScript, Tailwind CSS 4
- **AI Models:** 5 concurrent models with ensemble voting
- **Supported Formats:** Text, Images (PNG, JPG, GIF, WebP), Audio (MP3, WAV, M4A, OGG)
- **APIs:** ClaimBuster, Google Fact Check, DuckDuckGo Search
- **Deployment:** Docker support, CORS-enabled API, environment-based configuration